

An hourly kilometre-scale fine particulate matter reconstruction for Kandy, Sri Lanka: construction, validation without local ground truth, and measurement priorities

A. M. D. W. B. Alahakoon, Department of Environmental and Industrial Sciences, University of Peradeniya, Sri Lanka
2026

Undergraduate research project report (ENS4998), 2026. The report runs to 50,937 words, 34 figures and 13 tables, and rests on seventeen pre-registrations lodged before the analyses they cover.

The problem

Fine particulate matter (PM_{2.5}) carries the largest estimated burden of disease of any air pollutant, and monitoring is sparsest where concentrations are highest. Kandy is such a city: a valley settlement in the central highlands of Sri Lanka with 98,828 residents inside its municipal boundary (2012 census) and nearly 389,000 commuters entering on a typical weekday (World Bank, 2020), so that most of those exposed to its air during the day do not live in it. It has no reference-grade PM_{2.5} monitor in operation and two low-cost sensors with public records. Any PM_{2.5} field for Kandy is therefore a model output that almost no local observation can evaluate.

Aim

To reconstruct PM_{2.5} over the Kandy basin at hourly and one-kilometre resolution for 2019 to 2023 from data available everywhere; to establish how far each part of that reconstruction can be trusted; to determine the spatial scale below which it is no longer supported; and to identify the measurements that would most improve it.

The model

Concentration on the kilometre grid is a regional background, uniform across the city at any hour, plus a local increment redistributed by a pattern of unit spatial mean:

$$\text{PM}(x, y, t) = B(t) + \max(\Delta, 0) P(x, y, t) + \min(\Delta, 0) + \varepsilon(t) [P(x, y, t) - 1], \quad \Delta(t) = T(t) - B(t)$$

T , the basin-mean concentration, is a gradient-boosted model of the residual between the two local sensors and a composition-model prior, driven by reanalysis meteorology and satellite retrievals, shifted each year to a satellite-derived area mean and shaped to the sensors' daily and seasonal cycles. B is a regional-background estimate built from the same satellite level, the composition model and an air-mass classification, capped so that it never reaches the hourly total. P combines a satellite surface, a terrain-confinement term and a road-traffic emission surface carried through a terrain-steered dispersion step. The field averages to T (to within 0.56 per cent in practice), so misplacing material cannot create more of it, and each component declares which observations it may use. The local increment is a residual of this decomposition, not the mass emitted inside the city.

Validation without local ground truth

The two Kandy sensors trained and shaped T , so they cannot validate it. Three other lines of evidence are used, and their results are graded rather than combined. The complete model was built in ten monitored valley and basin cities with only two anchor stations, as at Kandy, and scored against withheld stations: the seasonal cycle transfers everywhere (correlation 0.943 to 0.996), the daily cycle unevenly (-0.325 to 0.977). The deployed temporal model was placed on a cross-city ladder of information budgets in 104 cities (exploratory, specified before scoring): where its anchor stations observed the period it cut daily error by 32.7 per cent against a generic sensorless estimate (correlation 0.864); where they did not, by 4.83 per cent, not distinguishable from zero. Its satellite-derived level sat 7.87 per cent above the withheld city means and its nominal 90 per cent interval covered a fraction 0.696. At Kandy an independent national record differs from the field by 0.7 and -2.6 per cent in two years, while three low-cost records sit below it.

The reconstructed field, and how far it can be trusted

The modelled annual basin mean is 17.1 to 21.0 micrograms per cubic metre over 2019 to 2023, above the World Health Organization guideline of 5 in every year; because the absolute level is unresolved, this describes the modelled field rather than an established exposure. The seasonal cycle is the best-supported property, with the maximum in the north-east monsoon (30.2 in December to February of 2023) and the minimum in the south-west monsoon (11.7). The day-to-day sequence is supported on days the sensors reported. Morning and evening peaks either side of a midday trough are robust in timing, but their depth depends on the humidity correction of the sensor record: with hourly rather than constant humidity the peak-to-trough ratio falls from 1.93 to 1.68 and the night falls below midday, while the level, the daily sequence and the local fraction barely move. The interval covers 72.4 per cent of sensor hours with misses almost all on one side, which points to a calibration bias at the sensors rather than proving the width right. Under the baseline decomposition the local increment is 0.483 of the modelled total, 0.433 to 0.492 across the tested ways of applying the cap and 0.547 with a 48-hour background window: a property of the decomposition, not a source apportionment.

What each additional observation is worth

Because every component declares its inputs, an observation can be withheld and the loss measured across monitored cities. A registered confirmation on 72 fresh cities found two stations used only to recalibrate a free estimate worth 8.53 per cent and a background series read daily worth 41.1 per cent. A later review showed the comparison was not like for like. In a post hoc re-analysis that used every source on the day, one station read daily reduced daily error by 42.1 per cent, two by 53.3 and five by 58.4, against about 11.0 per cent for stations used only as a calibration, and no ordering of station types was established. For Kandy the priority is continuously reporting stations whose readings reach the estimate every day, with one of reference grade to settle the level and the humidity correction.

Limitations

The field is not supported below the kilometre scale. Two sites three hundred metres apart inside one model cell differ in measured PM10 by a factor of 27.5, though on different days, so the pair illustrates sub-grid variation without isolating location; on densely monitored cities the spread inside a cell (1.22) exceeds the spread between cells (1.05). The dispersion step meant to place the local increment ranks held-out stations at 0.274, below the undispersed surface (0.371) and a free built-up layer (0.393), so it is a declared choice, not a validated improvement, and three to eight stations do not rank neighbourhoods usefully either. The map is a hypothesis for testing, not a ranking of neighbourhoods. The chemical check is a consistency check within a composition model, not a chemical measurement at Kandy.

Ongoing work and verification

The regulatory monitor at Kandy, once its record is released, will allow the level, the humidity correction and the daily shape to be checked against a reference instrument. A fresh pre-registered like-for-like test and an independent reproduction are the next steps for the cross-city results. Every value in the report is regenerated from source at build time by a gate that refuses to build when prose and data disagree (github.com/daminda1108/kandy_pm2.5_framework); twenty-three of 105 registered predictions were refuted and are reported as such.

Contact: 11daminda08@gmail.com · s20005@sci.pdn.ac.lk